

# Single Variable Regression

EC420 MSU Online

---

Justin Kirkpatrick  
Last updated May 28, 2021

## Lectures:

1. Population Regression Function: What are we after?
2. Ordinary Least Squares with an example
3. Goodness of fit measures
4. Interpretation of coefficients
5. Rescaling Y and X
6. Non-linear functional forms
7. Regression in R
8. Inference and hypothesis testing: Expectation of  $\hat{\beta}_1$
9. Inference and hypothesis testing: Variance of  $\hat{\beta}_1$
10. An example

# Population Regression Function

[top](#)

We have some data on two (or more, later) variables that we think move together in an interesting way.


- Wage and education
- Cigarette smoking and life expectancy
- COVID cases and vaccine rates

We want to quantify and test this relationship

- Predict a change
- Test a theory
- Win a bet?

We have a **sample**, but want to predict/test something about the population


Wage data used in Wooldridge `wage2`



Data from Blackburn and Neumark (1992), "Unobserved Ability, Efficiency Wages, and Interindustry Wage Differentials" *Quarterly Journal of Economics* 107, 1421-1436

The data looks like this:

wage	educ
769	12
808	18
825	14
650	12
562	11
1400	16



$N = 935$ ,  $\overline{wage} = 957.95$  and  $\overline{educ} = 13.47$

What we'd like to have is a function that tells us how *wage* and *educ* move together in the **population**

In a perfect world, we would have some function for  $X = educ$  and  $Y = wage$ :

$$g(x) = y$$

Where we give the function any realization of  $x$ , and it spits out exactly  $y$ .

But that isn't going to happen

Think about the data we just looked at - when  $educ = 12$  we observed  $wage = 769$  and  $wage = 650$ . The dream function doesn't exist! There are other things not accounted for besides  $educ$ .

So we settle for something that tells us about the **expectation** of  $Y$ .  
The Population Regression Function

$$E[Y|X] = \beta_0 + \beta_1 X$$

The *Population Regression Function* (PRF) describes the relationship between  $X$  and the **conditional expectation** of  $Y$ .

- $X$  and  $Y$  are random variables
- $\beta_0$  and  $\beta_1$  are **population parameters**
- We have restricted the  $E[Y|X]$  to be a *linear* function of  $X$ .
  - It can be drawn as a straight line with an intercept and constant slope
  - We will be estimating  $\beta_0$  and  $\beta_1$



The PRF:


$$E[Y|X] = \beta_0 + \beta_1 X$$

Let  $Y = \textit{wage}$  and  $X = \textit{educ}$


- $E[Y|X = x]$  gives us the expectation of  $Y$  (wage) conditional on some realized value of  $X = x$  (educ)
- So, if  $\textit{educ} = 16$ , then  $E[Y|X = 16] = \beta_0 + \beta_1 \times 16$ 
  - We can plug in any  $x_i$  and get the **expected value** of the paired  $y_i$

Question: Will the PRF return exactly  $y_i$  given a value  $x_i$ ?


No. It can only give us the expectation.




Ch. 2.1 of Wooldridge, example of a conceptual PRF. The line defines the PRF, the expectation of  $Y$  conditional on  $X$



This is the wage data. Each "blob" is an empirical histogram of the data for that value of *educ* (they are symmetrical). This is called a *violin plot*. It is the empirical counterpart of the [previous plot from Wooldridge](#)



Each point is the sample mean for each value of *educ*.



A (linear) PRF would be the straight line that best fits the data. **Regression fits that line.** A brief look at the line shows that it certainly won't be perfect!

What happens, then, if we want to write  $Y$  exactly?

The PRF gives us the *expectation* of  $Y$

- So we add a **stochastic error term**, the difference between  $E[Y|X]$  and  $Y$ :

$$Y = E[Y|X] + U = \beta_0 + \beta_1 X + U$$

This is the stochastic population regression function

$U$  is also the **population error term**, and is itself a **random variable**.

- It must be that  $E[U] = 0$

Now we can write our **simple linear regression model**:

$$y = \beta_0 + \beta_1 x + u$$

This is a statement about the relationship between observed realizations  $(y_i, x_i)$  based on the population parameters  $\beta_0, \beta_1$

We will call  $u$  the **error term** - it is the difference between the conditional expected mean and the observed  $y_i$  given a value of  $x_i$ .

- It might be different for two identical realizations of  $x_i$

Naturally, we would think that the "right" value of the population parameters,  $\beta = \{\beta_0, \beta_1\}^*$ , minimizes all of the  $u_i$  values in a sample. That's where Ordinary Least Squares comes in.

\* A parameter vector is just a list of numbers.



## The Sample Regression Function

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

## The "hats" are important

They mean we have a *sample estimate* of the population parameters.

- $\beta_0, \beta_1$  are the population
- $\hat{\beta}_0, \hat{\beta}_1$  are the sample estimates and will change when the sample changes
  - So they are random variables!

## Where did $u$ go?

Since we have a hat on  $y_i$ , there is no  $u$ , but  $\hat{y}_i \neq y_i$ .

- Define  $\hat{u}_i = \hat{y}_i - y_i$ .
- $\hat{u}_i$  is the *residual*.

To summarize:

The *PRF* is

$$E[Y|X] = \beta_0 + \beta_1 X$$

The simple linear regression model is:

$$y = \beta_0 + \beta_1 x + u$$

The SRF is:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

And if we want to write the sample regression model:

$$y_i = \hat{\beta}_0 + \hat{\beta}_1 x_i + \hat{u}_i$$

# Ordinary Least Squares

[top](#)

We have a linear PRF

$$E[Y|X] = \beta_0 + \beta_1 X$$

What happens, then, if we want to write  $Y$  exactly?

The PRF gives us the *expectation* of  $Y$

- So we add that **stochastic error term**, the difference between  $E[Y|X]$  and  $Y$ :

$$Y = E[Y|X] + U = \beta_0 + \beta_1 X + U$$

This is the stochastic population regression function

$U$  is also the **population error term**, and is itself a **random variable**.

- It must be that  $E[U] = 0$

And we can write our SRF:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

## How do we get those $\hat{\beta}$ 's in the SRF?

We make two assumptions:

**First**, if the expectation of  $Y$  equals  $\beta_0 + \beta_1 X$ , then *in expectation*,  $E[U] = 0$ .  
Because:

$$E[Y|X] = \beta_0 + \beta_1 X \quad \text{and} \quad Y = \beta_0 + \beta_1 X + U$$

**Second**, our first assumption should hold no matter what  $x$  is. So, it should be true that  $E[U|X] = 0$  for **all** possible values of  $X$ .

There are very important assumptions as they will define our Sample Regression Function (SRF).

Let's make these assumptions formal:

1.  $E[U] = 0$ .

- As long as there is a  $\beta_0$  (regardless of  $\beta_1$ ), this is true. We call this assumption **trivial**.

2.  $E[U|X] = E[U]$

- **Mean independence**. The **mean** of  $U$  is the same, regardless of the value of  $X$ :

These are **population moments**

- A **moment** is a specific attribute of a distribution

Economists spend a lot of time showing mean independence  $E[U|X] = E[U]$ .

Two quick math reminders before we introduce the Ordinary Least Squares (OLS) estimator for  $\beta$ :

$$\text{Cov}(Y, X) = E[YX] - E[Y]E[X]$$

and

$$\text{If } E[U] = 0$$

then

$$\text{Cov}(U, X) = E[UX] - E[U]E[X] = E[UX] - 0$$

And note that the simple linear regression model  $y = \beta_0 + \beta_1 x + u$  implies that:

$$u = y - \beta_0 - \beta_1 x$$

Since  $u = y - \beta_0 - \beta_1 x$ :

Let's write Assumption 1 and Assumption 2 using expectations of the **regression model from before**

- $E[U] = 0 \Rightarrow E[(y - \beta_0 - \beta_1 x)] = 0$
- $E[U|X] = 0 \Rightarrow E[x(y - \beta_0 - \beta_1 x)] = 0$ 
  - To see this, picture any expected value of  $x$ . Now, multiply it by 0.
- How many equations?
- How many unknowns?

## Let's solve for $\beta$ . To the "board"!

These are *moments*, and this way of deriving  $\beta$  is known as "method of moments".



What we just derived on the board depends on **population** moments:  $Cov(X, Y)$  and  $Var(X)$ .

But, just as before when we didn't know  $\mu$  but we could calculate  $\bar{y}$  (and we even know something about the distribution of  $\bar{y}$ ), we can calculate sample values for  $Cov(X, Y)$  and  $Var(X)$

Sample estimates get a  $\hat{\phantom{x}}$

So if we take a sample and calculate, from that sample,  $Var(X)$ , we would call it  $\widehat{Var}(X)$

First, let's tackle the estimate of  $\hat{\beta}_1$ .

We know how to calculate the sample covariance:

- $\widehat{Cov}(Y, X) = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$

We know how to calculate the sample variance:

- $\widehat{Var}(X) = \frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2$

$$\hat{\beta}_1 = \frac{\widehat{Cov}(Y, X)}{\widehat{Var}(X)} = \frac{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})}{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2}$$

What is important here is that **these are all observable in the data, and you know how to calculate them**. You know how to calculate  $\bar{x}$  and  $\bar{y}$ , you know how to sum things, and you know  $x_i$  and  $y_i$  in the data.


**As long as your assumptions hold**, you have an estimate of the PRF.

Second, let's tackle the *estimate* of  $\beta_0$ .

- We know, from the board, that  $\beta_0 = E[Y] - \beta_1 E[X]$
- We have a good, unbiased **sample** estimator for  $E[Y]$ :  $\bar{y}$ .
- And we have a good, unbiased **sample** estimator for  $E[X]$ :  $\bar{x}$ 
  - Plugging in:  $\bar{y} = \hat{\beta}_0 + \hat{\beta}_1 \bar{x}$

We don't observe  $\beta_0$ , but we can estimate it by taking sample mean of  $y$  and  $x$  (if we knew  $\hat{\beta}_1$ ).





The red line is the *sample regression function*, or *SRF*.

A couple important terms:

- The **fitted value**,  $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$
- The **residual**,  $\hat{u}_i = y_i - \hat{y}_i = y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i$

And note that:

- $y_i = \hat{\beta}_0 + \hat{\beta}_1 x_i + \hat{u}_i$ 
  - The  $\hat{u}_i$  "trues up" the fitted value.

Note that the residual is not the same as the error term.

- The residual is an empirical estimate from the sample
- The error term,  $u_i$ , is different

That error term is doing a lot of work here - it is covering *everything* that isn't  $x_1$ .  
So, what's inside the error term?


In  $u_i$

- Omitted variables
  - There might be another covariate,  $x_2$ , that is missing.
- Measurement error
  - That  $x$  might not be correctly measured.
- Non-linearities
  - Maybe there are some non-linear effects included in there.

These are all in  $u_i$ .


$$y_i = \beta_0 + \beta_1 x_1 + \underbrace{\beta_2 x_{omitted} + \beta_3 (x_1^* - x_1) + f(nonlinearities)}_{\text{other things, } u} + \tilde{u}_i$$

Our estimator,  $\hat{\beta}$  assumes allllll these things are 0 in expectation, no matter the value of  $x$



Wooldridge Fig. 2.4





Regression line for wage2 data

It will always be the case that, for any estimates  $\beta$  from a sample:

- $\sum_{i=1}^N (\hat{u}_i) = 0$
- $\sum_{i=1}^N (x_i \hat{u}_i) = 0$
- The point  $(\bar{y}, \bar{x})$  is always on the regression line

The "squares" part refers to the squaring of the error term.

The "least" part refers to a minimization of the (squared) error term.

Let's define the **sum of squared residuals** as:

$$SSR = \sum_{i=1}^N \hat{u}_i^2 = \sum_{i=1}^N (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2$$

And  $\beta$  is the "Least Squares" estimate if it minimizes  $SSR$ . How?

Take the derivative and set it equal to zero:

$$\frac{\partial SSR}{\partial \hat{\beta}_0} = 2 \sum (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) = 0$$

and

$$\frac{\partial SSR}{\partial \hat{\beta}_1} = 2 \sum (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) x_i = 0$$

## Terminology

$$y = \beta_0 + \beta_1 x + u$$

$y$  is called

- The dependent variable (DV)
- The "left hand side" (LHS)
- The outcome variable
- The response variable
- The target variable (in ML)

$x$  is called

- The independent variable
- The "right hand side" (RHS)
- The explanatory variable
- The control variable
- A covariate or a regressor

$u$  is called

- The residual (when  $\hat{u}$ )
- The error term (when  $u$ )

# Goodness of fit measures

[top](#)

## SSR, SSE, and SST

We know that  $\beta_{OLS}$  minimizes the sum of squares. How do we measure how good of a fit we get?

We previously defined  $SSR = \sum_{i=1}^N \hat{u}_i^2$

Define two more in addition to  $SSR$ :

- Sum of Squares Total:  $SST = \sum_{i=1}^N (y_i - \bar{y})^2$ 
  - $SST$  is a total sum of squares (notice no hats).
- Sum of Squares Explained:  $SSE = \sum_{i=1}^N (\hat{y}_i - \bar{y})^2$ 
  - $SSE$  can be thought of as how much is *explained* by  $\hat{y}_i$ , *relative to just guessing the obvious:  $\bar{y}$*

$$SST = SSR + SSE$$

The total variance is the sum of the variance of the residuals ("what isn't explained by your model") and the  $SSE$  ("the variance that is explained").

This is a *decomposition* of variance. It is important.

If your model fits well...

Then  $SSR = \sum_{i=1}^N \hat{u}_i^2$  is very small

- Because if your model fits well, then  $\hat{y}_i - y_i = \hat{u}_i$  is very small over all  $i$

Then  $SSE = \sum_{i=1}^N (\hat{y}_i - \bar{y})^2$  should be large

- Because if your model fits well, then you are explaining a lot of the observed deviations from  $\bar{y}$

## The $R^2$

$R^2$  ('r-square') is the comparison of  $SSE$  to  $SST$ . Since  $SSE < SST$  always, and both are always positive,  $0 < R^2 \leq 1$

$$R^2 = \frac{SSE}{SST} = 1 - \frac{SSR}{SST}$$


## Interpretation of $R^2$

The  $R^2$  is interpreted as the "fraction of variance in  $y$  explained by the model"

- Your regression, the SRF, is a model
- The variance being explained is the variance in the outcome,  $y$ .
- If it is 0, then  $SSE$  is zero. If  $SSE$  is zero, then  $y_i = \bar{y}$  for all  $i$ .
  - And in that case, your model isn't explaining *any* of the variance in  $y$ .



From earlier:




Wooldridge Fig. 2.4

# Interpretation of coefficients

[top](#)

Last time, we discussed a single variable regression from Wooldridge `wage2` where  $Y$  is *wage* and  $X$  is *educ*:

$$wage = \beta_0 + \beta_1 educ + u$$



This resulted in a  $\hat{\beta}_1 = 60.21$ . How do we interpret this coefficient?

Let's start with our simple linear regression model:

where *wage* and *educ* are random variables

$$wage = \beta_0 + \beta_1 educ + u$$

Our PRF is:

$$E[wage|educ] = \beta_0 + \beta_1 educ$$

"One additional year of education is associated with a  $\beta_1 = 60.21$  increase in *expected* monthly earnings, **all else held equal**"

- Why "expected"? We are estimating the PRF, so we are looking for the relationship between *expected* monthly earnings and education.
- Why "all else held equal"? **Because we have assumed that  $E[U|X] = 0$** , so our estimate tells us how  $E[Y]$  changes as  $X$  and not  $U$  changes.
  - $U$  is held at zero, no matter the  $X$

## Ceteris Paribus


is Latin for "all else held equal"

---

So the interpretation of  $\hat{\beta}_1$  is:

"The (estimated) increase in the expectation of *wage* associated with a 1-unit increase in *educ*, ceteris paribus"

The "all else held equal" part is very important.



- $\hat{\beta}_1$  is  $\frac{\Delta wage}{\Delta educ}$
- $\hat{\beta}_1$  is the slope of the line
  - The line is  $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$ , the *SRF*

## Regression in R

```
wage2 = wooldridge::wage2
myRegression = lm(wage ~ educ, data=wage2)
summary(myRegression)
```

```
##
## call:
## lm(formula = wage ~ educ, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -877.38 -268.63  -38.38   207.05  2148.26
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   146.952     77.715   1.891   0.0589 .
## educ          60.214      5.695  10.573 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 382.3 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

We see the  $\hat{\beta}_0$  labeled "intercept", and  $\hat{\beta}_1$ , the "coefficient on *educ*" is labeled with the variable name *educ*.

More on Regression in R [later](#)

# Rescaling Y and X

[top](#)




What happens if we re-scale the dependent variable, wage?

Maybe we have *wage* in dollars, but want it in thousands of dollars

We hope that it still gives us the same relationship

Define  $wage1000 = .001 \times wage$

- Any ideas what will happen to our coefficient?



Looks pretty similar, right? But the y-axis scale is very different.

A regression of:

$$wage1000 = \beta_0 + \beta_1 educ + u$$

```
##
## call:
## lm(formula = wage1000 ~ educ, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.87738 -0.26863 -0.03838  0.20705  2.14826
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.146952   0.077715   1.891   0.0589 .
## educ         0.060214   0.005695  10.573  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3823 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

$\hat{\beta}_1 = 0.06$  when we use *wage1000*

$\hat{\beta}_1 = 60.21$  when we use *wage*.

Re-scaling the dependent variable, *wage*, results in an equal rescaling of the coefficient.

The relationship predicted by the *SRF* stays the same.

Now, let's re-scale the *independent* variable

- That's the "right hand side" variable, *educ*.
- Let's do education in months:  $educMonths = educ \times 12$
- Any predictions on what will result?

```
##
## call:
## lm(formula = wage ~ educMonths, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -877.38 -268.63  -38.38   207.05  2148.26
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  146.9524    77.7150   1.891   0.0589 .
## educMonths     5.0179     0.4746  10.573  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 382.3 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

What was the result?

Re-scaling the independent variable simply rescales the coefficient by the *inverse* amount:

- $12 \times educ \Rightarrow \hat{\beta}_1^{new} = \frac{\hat{\beta}_1}{12}$

Re-scaling the dependent variable simply rescales the coefficient on it by an equal amount:

- $\hat{\beta}_1^{new} = \hat{\beta}_1 \times .001$

The relationship always remains the same

Let's take a look at the  $R^2$  of the original regression:

```
##
## call:
## lm(formula = wage ~ educ, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -877.38 -268.63  -38.38   207.05  2148.26
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   146.952     77.715   1.891   0.0589 .
## educ          60.214      5.695  10.573  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 382.3 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```



Now, the re-scaled dependent variable:

```
##
## Call:
## lm(formula = wage1000 ~ educ, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.87738 -0.26863 -0.03838  0.20705  2.14826
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.146952   0.077715   1.891   0.0589 .
## educ         0.060214   0.005695  10.573  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3823 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

And the re-scaled independent variable:

```
##
## Call:
## lm(formula = wage ~ educMonths, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -877.38 -268.63  -38.38   207.05  2148.26
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  146.9524    77.7150   1.891   0.0589 .
## educMonths    5.0179     0.4746  10.573  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 382.3 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

Heck, let's rescale both and look at the  $R^2$

```
##
## call:
## lm(formula = wage1000 ~ educMonths, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.87738 -0.26863 -0.03838  0.20705  2.14826
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.1469524  0.0777150   1.891  0.0589 .
## educMonths  0.0050179  0.0004746  10.573  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3823 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

The  $R^2$  is the same in every single one!

The "fraction of variance explained by the model" does not change.

Intuitively, you shouldn't be able to explain more variance simply by re-scaling a variable. The relationship that holds for wages and years of education must hold for 12 x years of education as well.

Since rescaling linearly doesn't matter, we can use a scale that is easiest to interpret and to read.

- *wage1000* in thousands of dollars is a lot easier to look at than the larger number we get using *wage*.
- You often don't want to have very extreme numbers of decimal places (e.g. a coefficient of .00000051 will be a lot easier to talk about if it's in millions: 5.1)

Now that we've seen an example, can we derive this result from the definition of  $\beta_1$ ?

$$\beta_1 = \frac{Cov(X, Y)}{Var(X)}$$

and if we rescale  $X$  by  $a$ :

$$\begin{aligned}\beta_1^{rescaled} &= \frac{Cov(aX, Y)}{Var(aX)} \\ &= \frac{aCov(X, Y)}{a^2Var(X)} \quad (\text{by rules of Cov and Var}) \\ &= \frac{Cov(X, Y)}{aVar(X)} \quad (\text{canceling}) \\ &= \frac{1}{a}\beta_1\end{aligned}$$

# Non-linear Functional Forms

[top](#)

What do we mean by "non-linear" function?


**A function** here is any mathematical operation or transformation that takes an input (usually called  $x$ ) and returns an output (usually called  $y$ ).

A non-linear function is any function where the graph is not a straight line.

- "Affine transformation" is the technical term for  $y = ax + b$ .
- "Non-affine transformation" is non-linear

Another way of thinking about non-linear functions is that  $\frac{\Delta y}{\Delta x}$  depends on the value of  $x$

- The slope of the graph changes as  $x$  changes.
- The slope at  $x_1$  (blue) is different than the slope at  $x_2$  (green)







In the previous slide, we saw a non-linear function, the exponential function,  $e^x$ . If we wanted a model to use in a regression that includes an exponential function, we could use:

$$y_i = \beta_0 + \beta_1 e^{x_i} + u_i$$

Note that the value of  $x_i$  is exponentiated.

- So this model has a non-linear term.
- It lets  $y$  respond to changes in  $x$  more flexibly

- but imposes that relationship whether it is appropriate (left) or not (right).




The most common non-linear transformation is the **polynomial**

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \cdots + u$$

For instance, plant growth rates over temperatures may be quadratic

- The *marginal effect* of an increase in temperature will be big and positive at lower temperatures.
- The *marginal effect* of an increase in temperature will be negative at very high temperatures.
- And somewhere in the middle, the *marginal effect* will be around zero.

The *marginal effect* is saying "the change in  $y$  per change in  $x$ ", or  $\frac{dy}{dx}$ .



If we have a polynomial relationship:

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + u$$

Then we can obtain the slope,  $\frac{dy}{dx}$  as the derivative of the relationship:

$$\frac{\partial y}{\partial x} = \beta_1 + 2\beta_2 x$$

If we propose a "higher order polynomial" relationship like:

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$$

Then we get a more complicated function for the slope at any  $x$ :

$$\frac{\partial y}{\partial x} = \beta_1 + 2\beta_2 x + 3\beta_3 x^2$$

There are other possible non-linear forms:  $\sqrt{x}$ , the natural log,  $\log_{10}$ , the inverse hyperbolic sine...

Even though these specifications are non-linear transformations, the regression is still **linear-in-parameters**

That is, all of the transformations we have discussed are still in the category of "linear models" because they are linear in the parameters.

So, our *PRF* (population regression function) is still linear, even with one of these transformations.

The quadratic specification,  $y = \beta_0 + \beta_1 x + \beta_2 x^2$  is particularly useful anytime you have an effect of  $x$  on  $y$  that dissipates or declines with increasing values of  $x$ .

Quick question: if the *effect* of  $x$  on  $y$  **declines** as  $x$  increases, then is the slope *increasing* or *decreasing* as  $x$  gets larger?

## An example:

In many cases, the effect of household income on some behavior may change as income increases.

- A low-income person may spend more on food when income increases
- But a high-income person may not spend much more on food when their income increases
  - But of course, the high-income will spend more on food than the low-income person.

We see these declining effects in many economic situations, but we also see increasing effects.

- Installing solar panels
- Others?

The quadratic "specification" can capture these phenomon.

## The natural log, $\ln(x)$

The natural log is the most common transformation. It is particularly useful because of the following:

$$\ln(1 + x) \approx x \quad \text{when} \quad x \approx 0$$

Let's say  $x^1 = x^0 + \Delta x$ .

$$\ln(x^1) - \ln(x^0) = \ln\left(\frac{x^1}{x^0}\right) = \ln\left(\frac{x^0 + \Delta x}{x^0}\right) = \ln\left(1 + \frac{\Delta x}{x^0}\right) \approx \frac{\Delta x}{x^0}$$

- This is the percent change in  $x$ :  $\frac{\Delta x}{x}$
- $100 \times [\ln(x^1) - \ln(x^0)] \approx \% \Delta x$



## The natural log, $\ln(x)$

Recall the formula for *elasticity*:  $\frac{\% \Delta y}{\% \Delta x} = \frac{\Delta y}{\Delta x} \times \frac{x}{y}$

And recall that, in a linear model (  $y = \beta_0 + \beta_1 x$  ), this elasticity is **not** constant:

$$\frac{\Delta y}{\Delta x} \times \frac{x}{y} = \beta_1 \times \frac{x}{y} = \beta_1 \times \frac{x}{\beta_0 + \beta_1 x + u}$$

But, when a model takes the form:  $\ln(y) = \beta_0 + \beta_1 \ln(x)$

$$\frac{\% \Delta y}{\% \Delta x} \approx \frac{\ln(y^1) - \ln(y^0)}{\ln(x^1) - \ln(x^0)} = \frac{\beta_1 [\ln(x^1) - \ln(x^0)]}{\ln(x^1) - \ln(x^0)} = \beta_1$$

The coefficient on a log-log model is the elasticity

$\ln(y) = \beta_0 + \beta_1 \ln(x)$  results in  $\beta_1$  being the elasticity of  $y$ , or "percent change in  $y$  from a 1 percent change in  $x$ ".

Econometrics is frequently about estimating that elasticity.

# Regression in R

top

## First, data

You should have already installed `wooldridge`. If not, type `install.packages('wooldridge')` directly in your console. Then, we can use R's built-in "data" function to load `wage2`

```
library(wooldridge)
wage2 = wooldridge::wage2 # creates a wage2 object
print(wage2[1:5,1:9]) # first 5 rows; first 9 columns
```

##		wage	hours	IQ	KWW	educ	exper	tenure	age	married
##	1	769	40	93	35	12	11	2	31	1
##	2	808	50	119	41	18	11	16	37	1
##	3	825	40	108	46	14	11	9	33	1
##	4	650	40	96	32	12	13	7	32	1
##	5	562	40	74	27	11	14	5	34	1

## Second, run the regression

We will use the `lm()` function. You will provide the regression formula and the name of the data to use.

The formula will be of the form  $y \sim x$ . You'll specify the data with `data = wage2`

```
MyRegression = lm(wage ~ educ, data=wage2)
print(MyRegression)
```

```
##
## call:
## lm(formula = wage ~ educ, data = wage2)
##
## Coefficients:
## (Intercept)      educ
##      146.95      60.21
```

Finally, we want a little more detail.

`MyRegression` is an R object. We can ask R to summarize it, and R will know to give us information about the regression:

```
summary(MyRegression)
```

```
##
## Call:
## lm(formula = wage ~ educ, data = wage2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -877.38 -268.63  -38.38   207.05  2148.26
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   146.952     77.715   1.891  0.0589 .
## educ          60.214      5.695  10.573 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 382.3 on 933 degrees of freedom
## Multiple R-squared:  0.107,    Adjusted R-squared:  0.106
## F-statistic: 111.8 on 1 and 933 DF,  p-value: < 2.2e-16
```

# Inference and hypothesis testing:

Expectation of  $\hat{\beta}_1$

[top](#)



We have a linear-in-parameters single-variable model:

$$y = \beta_0 + \beta_1 x + u$$

- "In terms of the random sample" (W2.5a):  $y_i = \beta_0 + \beta_1 x_i + u_i$
- "Fitting a line"
  - The PRF and the SRF
- $\hat{\beta}_1 = \frac{\widehat{Cov}(x,y)}{\widehat{Var}(x)}$
- $\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$
- **SST** (Sum of Squares Total) =  $\sum_{i=1}^N (y_i - \bar{y})^2$ 
  - **SSE** (Sum of Squares Explained) =  $\sum_{i=1}^N (\hat{y}_i - \bar{y})^2$
  - **SSR** (Sum of Squares Residual) =  $\sum_{i=1}^N (\hat{u}_i - \hat{\bar{u}})^2$

When we have a random variable with a population characteristic of interest

- For example,  $X$  with population mean  $\mu_X$

And a sample  $x_i$  of observed draws from the RV, then we can make a *hypothesis* about  $\mu_X$ :

- $H_0 : \mu_X = 0$  and  $H_A : \mu_X \neq 0$

Then, we can develop a sample *test statistic* for the population characteristic:

- $\bar{X} = \frac{1}{N} \sum x_i$

And we know two things about  $\bar{X}$ :

- $E[\bar{X}] = E[X] = \mu_X$
- $Var(\bar{X}) = \frac{\sigma_X^2}{N}$

If we're smart, we make a sample test statistic with a distribution that we know:

$$\frac{\bar{X} - H_0}{\sqrt{\frac{\hat{\sigma}^2}{N}}} \sim N(0, 1)$$

or if we don't know  $\sigma_X^2$

$$\frac{\bar{X} - H_0}{\sqrt{\frac{\hat{s}^2}{N}}} \sim t_{df}$$

We can test our hypothesis by comparing our sample test statistic result to the hypothesized value.

- If observed  $\bar{X} = 4$  and observed  $\frac{\hat{\sigma}_X}{\sqrt{N}} = 1$ , is  $H_0 : \mu_X = 0$  likely to be rejected?

So what if we want to test something about  $\beta_1$ ?

We can think of  $\beta_1$  as the test statistic for the relationship between  $x$  and  $y$

What do we need to test a hypothesis?

A **distribution**

- $E[\hat{\beta}_1]$
- $Var(\hat{\beta}_1)$
- $\hat{\beta}_1 \sim N(?, ?)$  (let's assume we know it's Normal for now)

If we did know these three things, we could test any interesting  $H_0$

- Anyone know one that might be interesting?

Now, remember that we are looking at  $\hat{\beta}$ , not  $\beta$  itself.

- $\beta$  is a population parameter,
  - It is unobserved
  - It is a constant
  - Because it is a constant, it can move in and out of **Expectations** and **Variances** as a constant would.
- $\hat{\beta}$  depends on the sample. It is therefore a random variable.
  - It has an expected value
  - It has a variance
  - We can use a statistical test on hypothesis about  $\hat{\beta}$ .

$\beta$  and  $\hat{\beta}$  are two different things, we are interested in whether or not they are the same in  $E$

## Gauss-Markov



Carl Friedrich Gauss



Andrey Markov

Both images courtesy of Wikimedia Commons

We will need to make the following four assumptions to get  $E[\hat{\beta}]$

## Gauss-Markov Assumptions

1. **SLR.1:** In the population,  $y$  is a linear function of the parameters,  $x$ , and  $u$ :  
$$y = \beta_0 + \beta_1 x + u$$
2. **SLR.2:** The sample  $(y_i, x_i) : i = 1, 2, \dots, n$  follows the population model and are independent.
3. **SLR.3:** "Sample Variation in the Explanatory (  $X$  ) Variable". That is,  $x_i$  is not the same for all  $i$ 's.
4. **SLR.4:** "Zero conditional mean".  $E[u|x] = 0$  for all  $x$ .

File these away for a minute. We'll need them.

## Expectation of the estimate: Bias

We know how to calculate, from our sample,  $\hat{\beta}$

We would hope (and will now prove) that  $E[\hat{\beta}] = \beta$

- This is the first step in deriving the distribution of  $\hat{\beta}$
- Section 2.5a of Wooldridge
  - If  $E[\hat{\beta}] = \beta$ , then the estimator is **unbiased**. Let's see if this is the case:

$$\hat{\beta}_1 = \frac{\widehat{Cov}(X, Y)}{\widehat{Var}(X)} = \frac{\frac{1}{N-1} \sum (x_i - \bar{x})(y_i - \bar{y})}{\frac{1}{N-1} \sum (x_i - \bar{x})^2} = \frac{\sum (x_i - \bar{x})y_i}{\sum (x_i - \bar{x})^2}$$

- The first equality is our derivation of  $\hat{\beta}_1$ .
- The second uses the definition of Covariance and Variance
- The third cancels out the  $\frac{1}{N-1}$  and does some simplification of the numerator (see Appendix A of Wooldridge)



Let's rewrite, then take expectations to see what the expectation of the estimate is:

$$\hat{\beta}_1 = \frac{\sum (x_i - \bar{x})y_i}{\sum (x_i - \bar{x})^2}$$

- Rewrite  $\sum (x_i - \bar{x})^2$  as  $SST_x$ . After all, it's the total sum of squared deviations from  $\bar{x}$ .
  - We are just adding that subscript to make sure we remember where it come from.
  - Remember, we originally introduced  $SST$  as the *Sum of Squares Total* in a regression and it referred to the total variance in  $Y$ , the left-hand-side (LHS) of our regression.
- Substitute our model for  $y_i$ :  $y_i = \beta_0 + \beta_1 x_i + u_i$
- Rename  $x_i - \bar{x}$  as  $d_i$ , for **d**eviations from  $\bar{x}$ .
  - This will make it easier to work with.

$$\hat{\beta}_1 = \frac{\sum (x_i - \bar{x})(\beta_0 + \beta_1 x_i + u_i)}{\sum (x_i - \bar{x})^2} = \frac{\sum (d_i \beta_0) + \sum (d_i \beta_1 x_i) + \sum (d_i u_i)}{SST_x}$$

Let's take a second and make sure everyone is on board here. Remember,  $d_i = x_i - \bar{x}$ .

Move the  $\beta$ 's out as they are constants:

$$\hat{\beta}_1 = \frac{\overbrace{\beta_0 \sum (d_i)}^{\text{First term}} + \overbrace{\beta_1 \sum (d_i x_i)}^{\text{Second term}} + \overbrace{\sum (d_i u_i)}^{\text{Third term}}}{SST_x}$$

In that numerator,  $\beta_0 \sum (d_i)$  must be 0 since  $\sum (x_i - \bar{x}) = 0$ . We can ignore it!

$$\hat{\beta}_1 = \frac{0}{SST_x} + \frac{\beta_1 \sum (d_i x_i)}{SST_x} + \frac{\sum (d_i u_i)}{SST_x}$$

The second term:

$$\frac{\beta_1 \sum(d_i x_i)}{SST_x} = \frac{\beta_1 \sum((x_i - \bar{x})x_i)}{SST_x} = \frac{\beta_1 \sum((x_i - \bar{x})(x_i - \bar{x}))}{SST_x} = \frac{\beta_1 SST_x}{SST_x}$$

And since  $SST_x$  is in the denominator and cancels, we will end up with  $\beta_1$ .

This is very important: notice that we now have the true value of beta in there.

$\beta_1$  is the **true beta**. It is *part of*  $\hat{\beta}_1$ , but there's still the third term:

$$\frac{\sum(d_i u_i)}{SST_x} = \frac{\sum((x_i - \bar{x})u_i)}{SST_x}$$
$$\hat{\beta}_1 = 0 + \beta_1 + \frac{\sum((x_i - \bar{x})u_i)}{SST_x}$$

We will say that the estimate of  $\beta_1$ ,  $\hat{\beta}_1$  is the true  $\beta$  plus some term.

$$\hat{\beta}_1 = \beta_1 + \frac{\sum((x_i - \bar{x})u_i)}{SST_x}$$

Conditional on the  $x_i$ 's (our sample), the entire source of randomness here is in  $u_i$ .

Now, we take the last step to show that the  $E[\hat{\beta}_1] = \beta_1$ .

We will need our four assumptions. Specifically, the fourth.

Our assumptions from before:

## Gauss-Markov Assumptions (fancy name for what you already know)

1. SLR.1: In the population,  $y$  is a linear function of the parameters,  $x$ , and  $u$ :  
$$y = \beta_0 + \beta_1 x + u$$
2. SLR.2: the sample  $(y_i, x_i) : i = 1, 2, \dots, n$  follows the population model and are independent.
3. SLR.3: "Sample Variation in the Explanatory (  $X$  ) Variable". That is,  $x_i$  is not the same for all  $i$ 's.
4. SLR.4: "Zero conditional mean".  $E[u|x] = 0$  for all  $x$ .

Now, we can go to our equation for  $\hat{\beta}_1$ :

$$\hat{\beta}_1 = \beta_1 + \frac{\sum((x_i - \bar{x})u_i)}{SST_x}$$

We can take  $E$  of each side:

$$E[\hat{\beta}_1] = E[\beta_1] + E\left[\frac{\sum((x_i - \bar{x})u_i)}{SST_x}\right]$$

$$E[\beta_1] = \beta_1.$$

For any value of  $x$ ,  $E[u|x] = 0$  under SLR.4.

- No matter what  $x$  or  $(x_i - \bar{x})$  is, once we condition on  $x$ , the second term is zero in expectation.

$$\Rightarrow E[\hat{\beta}_1] = \beta_1.$$

Our estimator,  $\hat{\beta}_1$  is **unbiased**, and we know it is distributed with mean of  $\beta_1$

$E[\hat{\beta}_0] = \beta_0$  is shown in Wooldridge 2.5a.

- "  $\hat{\beta}_0$  is an unbiased estimator of  $\beta_0$  "

Now, we simply need to fill in the variance of  $\hat{\beta}$  to have a test statistic for  $\beta$ .

# Inference and hypothesis testing:

Variance of  $\hat{\beta}_1$

[top](#)



## Gauss-Markov Assumptions

1. SLR.1: In the population,  $y$  is a linear function of the parameters,  $x$ , and  $u$ :  
$$y = \beta_0 + \beta_1 x + u$$
2. SLR.2: the sample  $(y_i, x_i) : i = 1, 2, \dots, n$  follows the population model and are independent.
3. SLR.3: "Sample Variation in the Explanatory (  $X$  ) Variable". That is,  $x_i$  is not the same for all  $i$ 's.
4. SLR.4: "Zero conditional mean".  $E[u|x] = 0$  for all  $x$ .

These get us to " $\hat{\beta}$  is unbiased"

## Gauss-Markov Assumptions

1. SLR.1: In the population,  $y$  is a linear function of the parameters,  $x$ , and  $u$ :  
$$y = \beta_0 + \beta_1 x + u$$
2. SLR.2: the sample  $(y_i, x_i) : i = 1, 2, \dots, n$  follows the population model and are independent.
3. SLR.3: "Sample Variation in the Explanatory (  $X$  ) Variable". That is,  $x_i$  is not the same for all  $i$ 's.
4. SLR.4: "Zero conditional mean".  $E[u|x] = 0$  for all  $x$ .

## Add one more assumption:

Add SLR.5:  $Var[u|x] = \sigma_u^2$  for all  $x$ .

- This is similar to the conditional mean, but says that every  $u_i$  is drawn from a variable whose distribution has the same value for  $\sigma^2$ .

SLR.5:  $Var[u|x] = \sigma_u^2$  for all  $x$

- This is similar to the conditional mean, but says that every  $u_i$  is drawn from a variable whose distribution has the same value for  $\sigma^2$ .
- We do **not** need this assumption to show that  $\hat{\beta}$  is an unbiased estimator for  $\beta$ 
  - But we do need this assumption to calculate the variance of  $\hat{\beta}$ .
- It does not mean that we know  $\sigma_u^2$ . We don't

Start with where we left off on  $\beta_1$ :

$$\hat{\beta}_1 = \beta_1 + \frac{\sum((x_i - \bar{x})u_i)}{SST_x}$$

Instead of taking the expectation as we did for proving unbiasedness, we take the **variance**:

$$Var(\hat{\beta}_1) = Var(\beta_1) + Var\left[\frac{\sum((x_i - \bar{x})u_i)}{SST_x}\right] + 2Cov\left(\beta_1, \left[\frac{\sum((x_i - \bar{x})u_i)}{SST_x}\right]\right)$$

- Because the variance of any constant (like  $\beta_1$ ) is 0, we can drop that 1st term.
- Because  $Cov(c, X) = 0$  when  $c$  is a constant, we can drop the  $2Cov(\dots)$  term.

This leaves us with:

$$Var(\hat{\beta}_1) = Var \left[ \frac{\sum((x_i - \bar{x})u_i)}{SST_x} \right] = Var \left[ \frac{1}{SST_x} \sum((x_i - \bar{x})u_i) \right]$$

We can condition on  $x_i$ 's again, and make the same argument that, conditional on  $x_i$ , we can take them out of the  $Var$  term.

- When we do this, we must **square** what we remove:

$$\begin{aligned} Var(\hat{\beta}_1) &= \frac{1}{SST_x^2} \times Var \left[ \sum(x_i - \bar{x})u_i \right] = \frac{1}{SST_x^2} \times \left[ \sum(x_i - \bar{x})^2 \right] Var(u_i) \\ &= \frac{SST_x}{SST_x^2} \sigma_u^2 = \frac{1}{SST_x} \sigma_u^2 \end{aligned}$$

So variance is:

$$Var(\hat{\beta}_1) = \frac{\sigma_u^2}{SST_x}$$

For any realization of  $\mathbf{x}$

- Variance of the estimator is increasing in  $\sigma_u^2$ .
- Variance of the estimator is decreasing in  $SST_x$ , variation in  $\mathbf{X}$ .

Good, but we don't know  $\sigma_u^2$ , do we?

- $\hat{u}$  seems like a good start.
- In our model,  $u_i$  is the *error*, but we observe  $\hat{u}$ , which is the *residual*.
  - $\hat{u}_i = u_i - (\hat{\beta}_0 - \beta_0) - (\hat{\beta}_1 - \beta_1)x_i$
  - So  $E[\hat{u}_i] = u_i$

As Wooldridge states: "the *error*,  $u$ , shows up in the equation containing the *population parameters*,  $\beta$ . The residual shows up in the *estimated* equation with  $\hat{\beta}$ ."

- Remember,  $u_i$  is not observed.
- But  $\hat{u}_i$  is observed.

We can use  $\sum_{i=1}^N \hat{u}_i^2$  as an estimator for  $\sigma_u^2$  if we make this small adjustment.

- $\hat{\sigma}_u^2 = \frac{1}{(N-2)} \sum_{i=1}^N \hat{u}_i^2 = \frac{SSR}{N-2}$
- This is because we know two things about  $\hat{u}$ :

$$\sum \hat{u} = 0$$

and

$$\sum x_i \hat{u}_i = 0$$

- We lose two **degrees of freedom**.
  - If we know all but two  $u_i$ 's, we could calculate the last two knowing these.
- **degrees of freedom** will be very important when we get to multiple regression.



This is the Standard Error of the Regression, SER

$$\hat{\sigma} = \sqrt{\hat{\sigma}^2} = \sqrt{\frac{\sum \hat{u}_i^2}{(N - 2)}}$$

We have used all five assumptions, but we can now say we know the distribution of  $\hat{\beta}$ :


$$\hat{\beta}_1 \sim N(\beta_1, \frac{\hat{\sigma}_u^2}{SST_x})$$

If we want to test a hypothesis about  $\beta_1$ , we now can.


But only **if** we assume homoskedasticity - that  $Var(u|x) = Var(u) = \sigma_u^2$ .

Let's take a look at this assumption briefly.

- Later on, we'll talk about how to adjust the Standard Error of the Regression for heteroskedasticity.



Homoskedasticity (from Wooldridge)



Heteroskedasticity (from Wooldridge)

# Single variable inference: an example

top

Let's work through an example with real data. Our goal is to take the data, calculate  $\beta_1$ ,  $\hat{se}(\hat{\beta}_1)$ , and test a hypothesis  $H_0 : \beta_1 = 0$ .

ID	Outcome	Dose
1	4.9	3
2	19.7	6
3	9.7	3
4	15.6	6
5	18.5	7

Statistic	Value
$\bar{y}$	13.68
$\bar{x}$	5
$SST_y = \sum (y_i - \bar{y})^2$	156.088
$SST_x = \sum (x_i - \bar{x})^2$	14
$\sum (y_i - \bar{y})(x_i - \bar{x})$	43.1

What is  $\hat{\beta}_1$ ?

What is  $\hat{\beta}_0$ ?

ID	Outcome	Dose	Fitted	Residual
1	4.9	3		
2	19.7	6		
3	9.7	3		
4	15.6	6		
5	18.5	7		

- Calculate  $\hat{y}$  using  $\beta_0$  and  $\beta_1$
- Calculate  $\hat{u}$  using  $y_i - \hat{y}$
- Calculate  $\hat{\sigma}_u^2$ 
  - Remember to divide by  $(n - 2)$  here for correct degrees of freedom

---

The formula for  $Var(\hat{\beta}_1)$  is  $\frac{\hat{\sigma}_u^2}{SST_x}$

- What is the distribution of  $\hat{\beta}_1$ ?

The formula for  $Var(\hat{\beta}_0)$  is  $\hat{\sigma}_u^2 \left[ \frac{1}{N} + \frac{\bar{x}^2}{SST_x} \right]$  (from Wooldridge)

- What is the distribution of  $\hat{\beta}_0$ ?

Using  $\hat{\beta}_1$  and the distribution of  $\hat{\beta}_1$ , what is the t-statistic under the null:  
 $H_0 : \beta_1 = 0$ :

$$\hat{t} = \frac{\hat{\beta}_1 - 0}{\hat{se}(\hat{\beta}_1)}$$

- Our  $\hat{t}$  is normally distributed\*, so we can check the p-value using the back of the Wooldridge book
  - Or, our "rule of thumb" of  $|\hat{t}| > 1.96$

\* Since  $s^2$  is estimated, this follows a t-distribution. However, for large enough N, the t distribution and the standard normal are very similar.



Check your work here:

```
##
## Call:
## lm(formula = Outcome ~ Dose, data = df)
##
## Residuals:
##      1      2      3      4      5
## -2.623  2.941  2.177 -1.159 -1.337
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.7129     3.9357  -0.435   0.6928
## Dose           3.0786     0.7464   4.124   0.0258 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.793 on 3 degrees of freedom
## Multiple R-squared:  0.8501,    Adjusted R-squared:  0.8001
## F-statistic: 17.01 on 1 and 3 DF,  p-value: 0.02584
```